

# BIN381\_Project\_Milestone 3\_MODELLING

Group F

2024-10-15

#Read and Split the Dataset

*#Load Libraries*

**library**(dplyr)

##

## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':

##

## filter, lag

## The following objects are masked from 'package:base':

##

## intersect, setdiff, setequal, union

**library**(rpart)

**library**(rpart.plot)

**library**(caret)

## Loading required package: ggplot2

## Loading required package: lattice

**library**(caTools)

**library**(randomForest)

## randomForest 4.7-1.2

## Type rfNews() to see new features/changes/bug fixes.

##

## Attaching package: 'randomForest'

## The following object is masked from 'package:ggplot2':

##

## margin

## The following object is masked from 'package:dplyr':

##

## combine

*#Load the dataset*

custData <- **read.csv**("CustData2\_Prepared.csv")

```
str(custData)
```

#Split the dataset to 80% training data and 20% testing data

```
set.seed(123)
```

## #Logistic Regression

```
logisticRegressionModel <- glm(formula = Eligible ~ . - Annual_Salary,  
                                data = train_data, family = 'binomial')
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
summary(logisticRegressionModel)
```

```
##
```

```
## Call:
```

```
## glm(formula = Eligible ~ . - Annual_Salary, family = "binomial",  
##      data = train_data)
```

```
##
```

```
## Coefficients: (7 not defined because of singularities)
```

```
##              Estimate Std. Error z value Pr(>|z|)  
## (Intercept) -4.234e+08  4.712e+11  -0.001  0.99928  
## Gross_Pay_Last_Paycheck 1.546e+00  6.487e-02  23.839 < 2e-16
```

```
***
```

```
## Gross_Year_To_Date -1.912e+00  1.674e+00  -1.142  0.25347  
## Gross_Year_To_Date_FRS_Contribution 5.049e+00  1.674e+00  3.017  0.00256
```

```
**
```

```
## Age 2.924e-02  1.174e-02  2.490  0.01278  
*
```

```
## Household_Size 3.623e-02  4.979e-02  0.728  0.46681
```

```
## Years_Residence -1.179e-02  1.545e-02  -0.763  0.44538
```

```
## Marital_Statusdivorced -1.326e-01  2.265e-01  -0.585  0.55833
```

```
## Marital_Statusmarried -2.633e-01  2.054e-01  -1.282  0.19997
```

```
## Marital_Statussingle -2.746e-01  2.048e-01  -1.341  0.18006
```

```
## Marital_Statuswidowed NA NA NA NA
```

```
## EducationBach. -1.487e-02  2.672e-02  -0.556  0.57790
```

```
## EducationHS.grad NA NA NA NA
```

```
## EducationMasters NA NA NA NA
```

```
## OccupationCleric. NA NA NA NA
```

```
## OccupationExec. NA NA NA NA
```

```
## OccupationProf. NA NA NA NA
```

```
## OccupationSales NA NA NA NA
```

```
## Salary_GroupLow 4.234e+08  4.712e+11  0.001  0.99928
```

```
## Salary_GroupMedium 4.234e+08  4.712e+11  0.001  0.99928
```

```
## Salary_GroupHigh 4.234e+08  4.712e+11  0.001  0.99928
```

```
## Salary_GroupVery.High 4.234e+08  4.712e+11  0.001  0.99928
```

```
## Frequency_Title 1.202e-04  3.372e-06  35.637 < 2e-16
```

```
***
```

```
## Frequency_Department -8.632e-06  1.416e-06  -6.095  1.1e-09
```

```
***
```

```
## Frequency_Country_ID 8.212e-07  5.429e-07  1.512  0.13042
```

```
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
##
```

```
## (Dispersion parameter for binomial family taken to be 1)
```

```
##
```

```
## Null deviance: 199391 on 153053 degrees of freedom
```

```
## Residual deviance: 42371 on 153036 degrees of freedom
```

```
## AIC: 42407
```

```
##
## Number of Fisher Scoring iterations: 20

#Make Predictions using Logistic Regression
logisticRegressionPrediction <- predict(logisticRegressionModel, newdata =
test_data, type = 'response')
head(logisticRegressionPrediction)

##           2           4           5           6           14
25
## 4.681861e-01 1.000000e+00 4.870882e-10 1.000000e+00 1.000000e+00
1.000000e+00

logisticRegressionY_pred = ifelse(logisticRegressionPrediction >0.5, 1, 0)

#Confusion matrix of Logistic Regression
logisticRegression_matrix <- table(actual = test_data$Eligible, predicted =
logisticRegressionY_pred)
logisticRegression_matrix

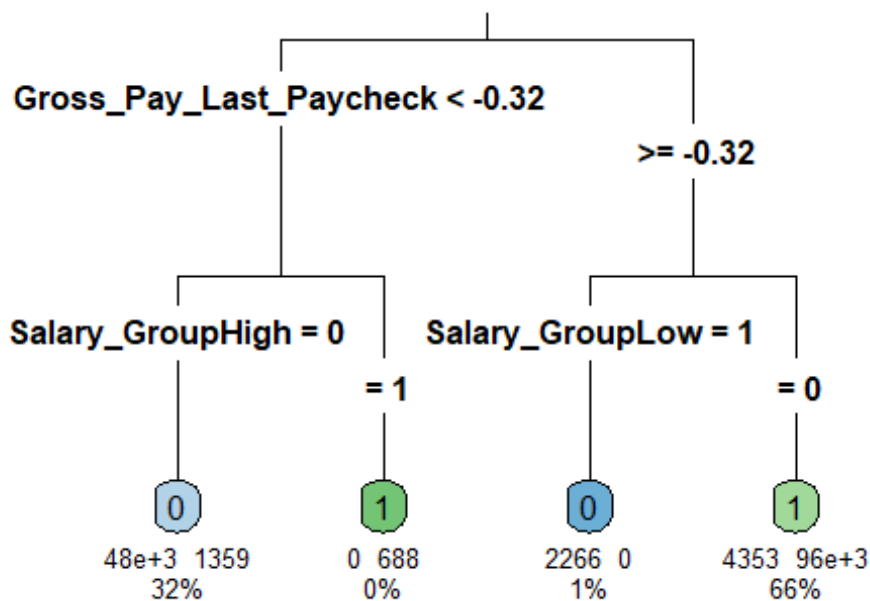
##      predicted
## actual      0      1
##      0 12483 1004
##      1   985 23791

logisticRegression_truePositive <- logisticRegression_matrix[1, 1]
logisticRegression_trueNegative <- logisticRegression_matrix[2, 2]
logisticRegression_falsePositive <- logisticRegression_matrix[1, 2]
logisticRegression_falseNegative <- logisticRegression_matrix[2, 1]

# Calculate Evaluation Metrics
logisticRegression_accuracy <- round((sum(diag(logisticRegression_matrix)) /
sum(logisticRegression_matrix)), 2)
logisticRegression_precision <- round(logisticRegression_truePositive /
(logisticRegression_truePositive + logisticRegression_falsePositive), 2)
logisticRegression_recall <- round(logisticRegression_truePositive /
(logisticRegression_truePositive + logisticRegression_falseNegative), 2)
logisticRegression_f1_score <- round(2 * (logisticRegression_precision *
logisticRegression_recall) / (logisticRegression_precision +
logisticRegression_recall), 2)

#Decision Tree

#Build Decision Tree Model
decisionTreeModel <- rpart(Eligible ~ . -Annual_Salary, data = train_data,
method = 'class')
```



```

# Make predictions on the test data
decisionTreePredictions <- predict(decisionTreeModel, newdata = test_data,
type = 'class')

#Confusion matrix
decisionTreeMatrix <- table(test_data$Eligible, decisionTreePredictions)
print(decisionTreeMatrix)

##    decisionTreePredictions
##          0          1
## 0 12450  1037
## 1   314 24462

decisionTreeTruePositive <- decisionTreeMatrix[1, 1]
decisionTreeTrueNegative <- decisionTreeMatrix[2, 2]
decisionTreeFalsePositive <- decisionTreeMatrix[1, 2]
decisionTreeFalseNegative <- decisionTreeMatrix[2, 1]


#Calculate Evaluation Metrics
decisionTreeAccuracy <- round((sum(diag(decisionTreeMatrix)) /
sum(decisionTreeMatrix)), 2)
decisionTreePrecision <- round(decisionTreeTruePositive /
(decisionTreeTruePositive + decisionTreeFalsePositive), 2)
decisionTreeRecall <- round(decisionTreeTruePositive /
(decisionTreeTruePositive + decisionTreeFalseNegative), 2)
decisionTreeF1Score <- round(2 * (decisionTreePrecision * decisionTreeRecall)
/ (decisionTreePrecision + decisionTreeRecall), 2)

```

```
#Reload the dataset
custData <- read.csv("CustData2_Prepared.csv")
custData$Eligible <- as.factor(custData$Eligible)

#Split the dataset to 80% training data and 20% testing data
set.seed(123)
train_index <- createDataPartition(custData$Eligible, p = 0.8, list = FALSE)
train_data <- custData[train_index, ]
test_data <- custData[-train_index, ]

#Build Random Forest Model
randomForest_model <- randomForest(Eligible ~ . -Annual_Salary, data =
train_data, ntree = 100, mtry = 3, importance = TRUE)
```



```
#Make Predictions Using Random Forest
randomForest_predictions <- predict(randomForest_model, newdata = test_data)

#Matrix for Random Forest
randomForest_cm <- confusionMatrix(as.factor(randomForest_predictions),
as.factor(test_data$Eligible))
randomForest_matrix <- randomForest_cm$table

randomForest_truePositive <- randomForest_matrix[1, 1]
randomForest_trueNegative <- randomForest_matrix[2, 2]
randomForest_falsePositive <- randomForest_matrix[1, 2]
```

```

randomForest_falseNegative <- randomForest_matrix[2, 1]

#Calculate Evaluation Metrics
randomForest_accuracy <- round((sum(diag(randomForest_matrix)) /
sum(randomForest_matrix)), 2)
randomForest_precision <- round(randomForest_truePositive /
(randomForest_truePositive + randomForest_falsePositive), 2)
randomForest_recall <- round(randomForest_truePositive /
(randomForest_truePositive + randomForest_falseNegative), 2)
randomForest_f1_score <- round(2 * (randomForest_precision *
randomForest_recall) / (randomForest_precision + randomForest_recall), 2)

```

#Model Evaluation

```

##Print Evaluation Metrics of All Models
cat("Logistic Regression Accuracy:", logisticRegression_accuracy, "\n")
## Logistic Regression Accuracy: 0.95

cat("Logistic Regression Precision:", logisticRegression_precision, "\n")
## Logistic Regression Precision: 0.93

cat("Logistic Regression Recall:", logisticRegression_recall, "\n")
## Logistic Regression Recall: 0.93

cat("Logistic Regression F1-score:", logisticRegression_f1_score, "\n")
## Logistic Regression F1-score: 0.93

cat("Decision Tree Accuracy:", decisionTreeAccuracy, "\n")
## Decision Tree Accuracy: 0.96

cat("Decision Tree Precision:", decisionTreePrecision, "\n")
## Decision Tree Precision: 0.92

cat("Decision Tree Recall:", decisionTreeRecall, "\n")
## Decision Tree Recall: 0.98

cat("Decision Tree F1-score:", decisionTreeF1Score, "\n")
## Decision Tree F1-score: 0.95

cat("Random Forest Accuracy:", randomForest_accuracy, "\n")
## Random Forest Accuracy: 0.97

cat("Random Forest Precision:", randomForest_precision, "\n")
## Random Forest Precision: 0.98

```

```
cat("Random Forest Recall:", randomForest_recall, "\n")
## Random Forest Recall: 0.93
cat("Random Forest F1-score:", randomForest_f1_score, "\n")
## Random Forest F1-score: 0.95

#Save Models

#Save the logistic regression model
saveRDS(logisticRegressionModel, file = "logistic_regression_model.rds")

#Save the decision tree model
saveRDS(decisionTreeModel, file = "decision_tree_model.rds")

#Save the random forest model
saveRDS(randomForest_model, file = "random_forest_model.rds")
```